



The Ultimate Must-Read AI Research Paper List

For Serious AI/ML/DL/NLP/CV/GAN/LLM Practitioners

1 ABSOLUTE MUST-KNOWS (Start Here - These 15 Papers Define Modern AI)

1. **Attention Is All You Need** | Vaswani et al. | 2017 [\[1\]](#)

- **Core Foundation:** Introduced Transformer architecture, powering all modern LLMs
- **Citations:** 173,000+ | **Impact:** Revolutionary - enabled ChatGPT, GPT-4, BERT
- **Link:** [arXiv](#) | [Code](#)

2. **ImageNet Classification with Deep Convolutional Neural Networks (AlexNet)** | Krizhevsky, Sutskever, Hinton | 2012 [\[1\]](#)

- **Core Foundation:** Started the modern deep learning revolution
- **Citations:** 100,000+ | **Impact:** Made deep learning famous, sparked AI boom
- **Link:** [Paper](#) | [Code](#)

3. **Generative Adversarial Networks** | Ian Goodfellow et al. | 2014 [\[1\]](#)

- **Core Foundation:** Introduced GAN framework for generative modeling
- **Citations:** 60,000+ | **Impact:** Revolutionized image generation, synthetic data
- **Link:** [arXiv](#) | [Code](#)

4. **Deep Residual Learning for Image Recognition (ResNet)** | He et al. | 2015 [\[1\]](#)

- **Core Foundation:** Solved vanishing gradient problem with residual connections
- **Citations:** 200,000+ | **Impact:** Enabled training very deep networks
- **Link:** [arXiv](#) | [Code](#)

5. **BERT: Pre-training of Deep Bidirectional Transformers** | Devlin et al. | 2018 [\[1\]](#)

- **Core Foundation:** Bidirectional language understanding, pre-training paradigm
- **Citations:** 80,000+ | **Impact:** Transformed NLP, established transfer learning
- **Link:** [arXiv](#) | [Code](#)

6. **Human-level control through deep reinforcement learning (DQN)** | Mnih et al. | 2015 [\[2\]](#)

- **Core Foundation:** Combined deep learning with reinforcement learning

- **Citations:** 37,000+ | **Impact:** Achieved human-level game performance
- **Link:** [Nature](#) | [Code](#)

7. Adam: A Method for Stochastic Optimization | Kingma & Ba | 2014^[3]

- **Core Foundation:** Most widely used optimization algorithm in deep learning
- **Citations:** 226,000+ | **Impact:** Default optimizer for neural networks
- **Link:** [arXiv](#) | [Code](#)

8. Playing Atari with Deep Reinforcement Learning | Mnih et al. | 2013^[4]

- **Core Foundation:** First successful deep RL, introduced DQN concept
- **Citations:** 15,000+ | **Impact:** Proved RL could work with raw pixels
- **Link:** [arXiv](#) | [Code](#)

9. Neural Machine Translation by Jointly Learning to Align and Translate | Bahdanau et al. | 2014^[1]

- **Core Foundation:** Introduced attention mechanism (pre-Transformer)
- **Citations:** 25,000+ | **Impact:** Enabled modern NMT and attention-based models
- **Link:** [arXiv](#) | [Code](#)

10. Very Deep Convolutional Networks for Large-Scale Image Recognition (VGG) | Simonyan & Zisserman | 2014^[5]

- **Core Foundation:** Established deeper networks with small filters
- **Citations:** 80,000+ | **Impact:** Standard architecture for many CV applications
- **Link:** [arXiv](#) | [Code](#)

11. Long Short-Term Memory | Hochreiter & Schmidhuber | 1997^[6]

- **Core Foundation:** Solved vanishing gradients in RNNs, enabled sequence modeling
- **Citations:** 100,000+ | **Impact:** Foundation for pre-Transformer NLP
- **Link:** [Paper](#) | [Code](#)

12. Learning representations by back-propagating errors | Rumelhart, Hinton, Williams | 1986^[6]

- **Core Foundation:** Backpropagation algorithm that enables neural network training
- **Citations:** 80,000+ | **Impact:** Made modern neural networks possible
- **Link:** [Paper](#) | Essential theory

13. Auto-Encoding Variational Bayes (VAE) | Kingma & Welling | 2013^[1]

- **Core Foundation:** Principled approach to generative modeling with latent variables
- **Citations:** 30,000+ | **Impact:** Foundation for many generative models
- **Link:** [arXiv](#) | [Code](#)

14. Language Models are Few-Shot Learners (GPT-3) | Brown et al. | 2020^[7]

- **Core Foundation:** Demonstrated emergent abilities through scaling
- **Citations:** 25,000+ | **Impact:** Sparked current LLM revolution
- **Link:** [arXiv](#) | [API](#)

15. Scaling Laws for Neural Language Models | Kaplan et al. | 2020^[7]

- **Core Foundation:** Mathematical relationship between scale and performance
- **Citations:** 5,000+ | **Impact:** Guides modern LLM development
- **Link:** [arXiv](#) | Theory foundation

II FOUNDATIONAL THEORY (*Essential Mathematical & Theoretical Foundations*)

16. A Logical Calculus of the Ideas Immanent in Nervous Activity | McCulloch & Pitts | 1943^[8]

- **Historical Foundation:** First mathematical model of neural networks
- **Impact:** Established neural computation as viable concept
- **Link:** [Paper](#)

17. The Perceptron: A Probabilistic Model | Rosenblatt | 1958^[8]

- **Historical Foundation:** First learning algorithm for neural networks
- **Impact:** Established concept of machine learning from data
- **Link:** [Paper](#)

18. Computing Machinery and Intelligence | Turing | 1950^[1]

- **Historical Foundation:** The Turing Test, philosophical foundations of AI
- **Impact:** Defined what it means for machines to think
- **Link:** [Paper](#)

19. Support Vector Networks | Cortes & Vapnik | 1995^[1]

- **Core Theory:** Foundation of kernel methods and modern ML theory
- **Citations:** 80,000+ | **Impact:** Pre-deep learning ML workhorse
- **Link:** [Paper](#)

20. Random Forests | Breiman | 2001^[9]

- **Core Theory:** Ensemble methods and feature importance
- **Citations:** 100,000+ | **Impact:** Still widely used for tabular data
- **Link:** [Paper](#)

DEEP LEARNING FOUNDATIONS (*Papers That Enabled Deep Learning*)

21. Deep Learning | LeCun, Bengio, Hinton | 2015^[1]

- **Review Paper:** Comprehensive overview by the "Godfathers of AI"
- **Citations:** 50,000+ | **Impact:** Definitive introduction to deep learning
- **Link:** [Nature](#)

22. Gradient-Based Learning Applied to Document Recognition (LeNet) | LeCun et al. | 1998^[1]

- **Foundation:** First successful CNN, introduced convolutional layers
- **Citations:** 35,000+ | **Impact:** Blueprint for modern CNNs
- **Link:** [Paper](#)

23. A Fast Learning Algorithm for Deep Belief Nets | Hinton et al. | 2006^[1]

- **Renaissance:** Revived deep learning after AI winter
- **Citations:** 25,000+ | **Impact:** Started modern deep learning era
- **Link:** [Paper](#)

24. Reducing the Dimensionality of Data with Neural Networks | Hinton & Salakhutdinov | 2006^[1]

- **Foundation:** Showed deep autoencoders for dimensionality reduction
- **Citations:** 20,000+ | **Impact:** Demonstrated power of deep architectures
- **Link:** [Science](#)

25. Understanding the difficulty of training deep feedforward neural networks | Glorot & Bengio | 2010^[9]

- **Theory:** Weight initialization and training dynamics
- **Citations:** 15,000+ | **Impact:** Xavier initialization still used today
- **Link:** [Paper](#)

COMPUTER VISION MILESTONES (*Visual Intelligence Breakthroughs*)

26. You Only Look Once: Unified, Real-Time Object Detection (YOLO) | Redmon et al. | 2015^[5]

- **Revolution:** Real-time object detection in single pass
- **Citations:** 40,000+ | **Impact:** Enabled real-time computer vision
- **Link:** [arXiv](#) | [Code](#)

27. Rich feature hierarchies for accurate object detection (R-CNN) | Girshick et al. | 2014^[5]

- **Foundation:** Region-based object detection
- **Citations:** 25,000+ | **Impact:** Started modern object detection

- **Link:** [arXiv](#) | [Code](#)

28. Faster R-CNN: Towards Real-Time Object Detection | Ren et al. | 2015^[5]

- **Evolution:** Unified detection pipeline with Region Proposal Networks
- **Citations:** 35,000+ | **Impact:** Standard for two-stage detection
- **Link:** [arXiv](#) | [Code](#)

29. U-Net: Convolutional Networks for Biomedical Image Segmentation | Ronneberger et al. | 2015^[5]

- **Architecture:** Skip connections for dense prediction tasks
- **Citations:** 60,000+ | **Impact:** Standard for segmentation tasks
- **Link:** [arXiv](#) | [Code](#)

30. An Image is Worth 16×16 Words: Transformers for Image Recognition (ViT) | Dosovitskiy et al. | 2020^[10]

- **Revolution:** Applied Transformers to computer vision
- **Citations:** 15,000+ | **Impact:** Started vision transformer era
- **Link:** [arXiv](#) | [Code](#)

31. Feature Pyramid Networks for Object Detection | Lin et al. | 2016^[5]

- **Architecture:** Multi-scale feature representation
- **Citations:** 15,000+ | **Impact:** Improved detection across scales
- **Link:** [arXiv](#) | [Code](#)

32. Mask R-CNN | He et al. | 2017^[5]

- **Extension:** Instance segmentation framework
- **Citations:** 20,000+ | **Impact:** State-of-the-art instance segmentation
- **Link:** [arXiv](#) | [Code](#)

II NATURAL LANGUAGE PROCESSING EVOLUTION (*Language Understanding Journey*)

33. Distributed Representations of Words and Phrases (Word2Vec) | Mikolov et al. | 2013^[11]

- **Foundation:** Dense word embeddings, semantic relationships
- **Citations:** 50,000+ | **Impact:** Modern NLP foundation
- **Link:** [arXiv](#) | [Code](#)

34. Sequence to Sequence Learning with Neural Networks | Sutskever et al. | 2014^[11]

- **Architecture:** Encoder-decoder for sequence transduction
- **Citations:** 25,000+ | **Impact:** Foundation for machine translation

- **Link:** [arXiv](#) | [Code](#)

35. GloVe: Global Vectors for Word Representation | Pennington et al. | 2014 [\[1\]](#)

- **Method:** Global statistics for word embeddings
- **Citations:** 30,000+ | **Impact:** Alternative to Word2Vec
- **Link:** [Paper](#) | [Code](#)

36. ELMo: Deep contextualized word representations | Peters et al. | 2018 [\[1\]](#)

- **Innovation:** Context-dependent embeddings
- **Citations:** 12,000+ | **Impact:** Bridged to modern contextual models
- **Link:** [arXiv](#) | [Code](#)

37. Improving Language Understanding by Generative Pre-Training (GPT-1) | Radford et al. | 2018 [\[1\]](#)

- **Paradigm:** Unsupervised pre-training for language understanding
- **Citations:** 8,000+ | **Impact:** Started GPT lineage
- **Link:** [Paper](#) | [Code](#)

38. Language Models are Unsupervised Multitask Learners (GPT-2) | Radford et al. | 2019 [\[1\]](#)

- **Scaling:** Larger model, emergent capabilities
- **Citations:** 12,000+ | **Impact:** Showed scaling potential
- **Link:** [Paper](#) | [Code](#)

39. RoBERTa: A Robustly Optimized BERT Pretraining Approach | Liu et al. | 2019 [\[1\]](#)

- **Optimization:** Improved BERT training recipe
- **Citations:** 8,000+ | **Impact:** Better BERT-style models
- **Link:** [arXiv](#) | [Code](#)

□ GENERATIVE MODELS (*Synthesis & Creation*)

40. Unsupervised Representation Learning with Deep Convolutional GANs (DCGAN) |

Radford et al. | 2015 [\[11\]](#)

- **Architecture:** Stabilized GAN training with convolutions
- **Citations:** 15,000+ | **Impact:** Made GANs practical
- **Link:** [arXiv](#) | [Code](#)

41. Conditional Generative Adversarial Nets | Mirza & Osindero | 2014 [\[1\]](#)

- **Extension:** Controlled generation with conditioning
- **Citations:** 12,000+ | **Impact:** Enabled controllable synthesis
- **Link:** [arXiv](#) | [Code](#)

42. Wasserstein GAN | Arjovsky et al. | 2017^[11]

- **Theory:** Improved training stability with Earth Mover's Distance
- **Citations:** 8,000+ | **Impact:** More stable GAN training
- **Link:** [arXiv](#) | [Code](#)

43. A Style-Based Generator Architecture (StyleGAN) | Karras et al. | 2018^[11]

- **Innovation:** Disentangled style control in generation
- **Citations:** 6,000+ | **Impact:** High-quality controllable faces
- **Link:** [arXiv](#) | [Code](#)

44. Denoising Diffusion Probabilistic Models | Ho et al. | 2020^[17]

- **New Paradigm:** Diffusion process for generation
- **Citations:** 8,000+ | **Impact:** Alternative to GANs, enabled DALL-E 2
- **Link:** [arXiv](#) | [Code](#)

45. Progressive Growing of GANs | Karras et al. | 2017^[11]

- **Training:** Progressive resolution for stable training
- **Citations:** 4,000+ | **Impact:** High-resolution image generation
- **Link:** [arXiv](#) | [Code](#)

46. CycleGAN: Unpaired Image-to-Image Translation | Zhu et al. | 2017^[11]

- **Innovation:** Translation without paired data
- **Citations:** 15,000+ | **Impact:** Artistic style transfer, domain adaptation
- **Link:** [arXiv](#) | [Code](#)

REINFORCEMENT LEARNING (*Learning Through Interaction*)

47. Asynchronous Methods for Deep Reinforcement Learning (A3C) | Mnih et al. | 2016^[4]

- **Architecture:** Asynchronous actor-critic methods
- **Citations:** 8,000+ | **Impact:** Scalable policy gradient methods
- **Link:** [arXiv](#) | [Code](#)

48. Proximal Policy Optimization (PPO) | Schulman et al. | 2017^[4]

- **Algorithm:** Simple, effective policy gradient method
- **Citations:** 5,000+ | **Impact:** Most popular RL algorithm today
- **Link:** [arXiv](#) | [Code](#)

49. Trust Region Policy Optimization (TRPO) | Schulman et al. | 2015^[4]

- **Theory:** Theoretical guarantees for policy updates

- **Citations:** 4,000+ | **Impact:** Principled policy improvement
- **Link:** [arXiv](#) | [Code](#)

50. Deep Deterministic Policy Gradient (DDPG) | Lillicrap et al. | 2015^[4]

- **Method:** Actor-critic for continuous control
- **Citations:** 8,000+ | **Impact:** Continuous action spaces
- **Link:** [arXiv](#) | [Code](#)

51. Mastering the Game of Go with Deep Neural Networks (AlphaGo) | Silver et al. | 2016^[1]

- **Breakthrough:** Superhuman performance in Go
- **Citations:** 12,000+ | **Impact:** Proved RL potential in complex games
- **Link:** [Nature](#) | [Code](#)

52. Rainbow: Combining Improvements in Deep Reinforcement Learning | Hessel et al. | 2017^[4]

- **Integration:** Combined multiple DQN improvements
- **Citations:** 3,000+ | **Impact:** State-of-the-art value-based RL
- **Link:** [arXiv](#) | [Code](#)

❖ OPTIMIZATION & TRAINING (*Making Learning Work*)

53. Batch Normalization: Accelerating Deep Network Training | Ioffe & Szegedy | 2015^[9]

- **Technique:** Normalizing layer inputs for stable training
- **Citations:** 40,000+ | **Impact:** Standard in most architectures
- **Link:** [arXiv](#) | [Code](#)

54. Dropout: A Simple Way to Prevent Neural Networks from Overfitting | Srivastava et al. | 2014^[1]

- **Regularization:** Random neuron dropping during training
- **Citations:** 50,000+ | **Impact:** Standard regularization technique
- **Link:** [Paper](#) | [Code](#)

55. Layer Normalization | Ba et al. | 2016^[9]

- **Alternative:** Layer-wise normalization for RNNs/Transformers
- **Citations:** 8,000+ | **Impact:** Essential for Transformer training
- **Link:** [arXiv](#) | [Code](#)

56. On the Convergence of Adam and Beyond | Reddi et al. | 2019^[12]

- **Theory:** Fixed Adam convergence issues
- **Citations:** 2,000+ | **Impact:** AMSGrad and improved optimizers

- **Link:** [arXiv](#) | [Code](#)

□ MODERN LLMs & FOUNDATION MODELS (*Current AI Revolution*)

57. LLaMA: Open and Efficient Foundation Language Models | Touvron et al. | 2023 [\[7\]](#)

- **Open Source:** Competitive open-source LLM
- **Citations:** 8,500+ | **Impact:** Democratized access to powerful LLMs
- **Link:** [arXiv](#) | [Code](#)

58. GPT-4 Technical Report | OpenAI | 2023 [\[7\]](#)

- **Milestone:** Multimodal large language model
- **Citations:** 3,400+ | **Impact:** Current state-of-the-art general AI
- **Link:** [arXiv](#) | [API](#)

59. PaLM: Scaling Language Modeling with Pathways | Chowdhery et al. | 2022 [\[7\]](#)

- **Scale:** 540B parameter model with emergent abilities
- **Citations:** 2,000+ | **Impact:** Demonstrated benefits of extreme scale
- **Link:** [arXiv](#)

60. InstructGPT: Training language models to follow instructions | Ouyang et al. | 2022 [\[7\]](#)

- **Alignment:** Human feedback for better instruction following
- **Citations:** 4,000+ | **Impact:** Foundation for ChatGPT
- **Link:** [arXiv](#)

61. Constitutional AI: Harmlessness from AI Feedback | Bai et al. | 2022 [\[7\]](#)

- **Safety:** AI system training itself to be less harmful
- **Citations:** 1,500+ | **Impact:** Scalable AI alignment approach
- **Link:** [arXiv](#)

62. Flamingo: a Visual Language Model for Few-Shot Learning | Alayrac et al. | 2022 [\[7\]](#)

- **Multimodal:** Few-shot learning with vision and language
- **Citations:** 2,000+ | **Impact:** Advanced multimodal reasoning
- **Link:** [arXiv](#)

□ MULTIMODAL & CROSS-MODAL (*Beyond Single Modalities*)

63. CLIP: Learning Transferable Visual Representations from Natural Language | Radford et al. | 2021 [\[10\]](#)

- **Innovation:** Joint vision-language representation learning
- **Citations:** 8,000+ | **Impact:** Foundation for multimodal AI

- **Link:** [arXiv](#) | [Code](#)

64. DALL-E: Creating Images from Text | Ramesh et al. | 2021 [\[1\]](#)

- **Breakthrough:** High-quality text-to-image generation
- **Citations:** 3,000+ | **Impact:** Demonstrated multimodal generation
- **Link:** [arXiv](#) | [API](#)

65. DALL-E 2: Hierarchical Text-Conditional Image Generation | Ramesh et al. | 2022 [\[2\]](#)

- **Evolution:** Improved text-to-image with diffusion
- **Citations:** 2,000+ | **Impact:** Mainstream AI art generation
- **Link:** [arXiv](#) | [API](#)

66. Segment Anything | Kirillov et al. | 2023 [\[3\]](#)

- **Foundation:** Universal image segmentation model
- **Citations:** 5,300+ | **Impact:** Zero-shot segmentation capabilities
- **Link:** [arXiv](#) | [Code](#)

67. LLaVA: Visual Instruction Tuning | Liu et al. | 2023 [\[4\]](#)

- **Innovation:** Large multimodal model for visual chat
- **Citations:** 2,800+ | **Impact:** Multimodal conversational AI
- **Link:** [arXiv](#) | [Code](#)

□ THEORETICAL BREAKTHROUGHS (*Understanding Why Things Work*)

68. The Lottery Ticket Hypothesis | Frankle & Carbin | 2018 [\[5\]](#)

- **Theory:** Sparse subnetworks can match full network performance
- **Citations:** 3,000+ | **Impact:** Network pruning and efficiency
- **Link:** [arXiv](#) | [Code](#)

69. Deep Double Descent | Nakkiran et al. | 2019 [\[6\]](#)

- **Phenomenon:** Performance improves again after classical overfitting
- **Citations:** 1,500+ | **Impact:** Rethinking bias-variance tradeoff
- **Link:** [arXiv](#)

70. Neural Tangent Kernel | Jacot et al. | 2018 [\[7\]](#)

- **Theory:** Infinite-width neural networks as kernel methods
- **Citations:** 2,500+ | **Impact:** Theoretical understanding of training
- **Link:** [arXiv](#) | [Code](#)

71. Understanding deep learning requires rethinking generalization | Zhang et al. | 2016 [\[8\]](#)

- **Theory:** Deep networks can memorize random labels
- **Citations:** 8,000+ | **Impact:** Challenged traditional ML theory
- **Link:** [arXiv](#)

72. **Grokking: Generalization Beyond Overfitting** | Power et al. | 2022^[10]

- **Phenomenon:** Sudden generalization long after overfitting
- **Citations:** 500+ | **Impact:** New understanding of learning dynamics
- **Link:** [arXiv](#)

RECENT BREAKTHROUGHS (2022-2025) (*Cutting-Edge Developments*)

73. **Mamba: Linear-Time Sequence Modeling with Selective State Spaces** | Gu & Dao | 2023^[13]

- **Architecture:** Alternative to Transformers with linear complexity
- **Citations:** 2,000+ | **Impact:** Potential Transformer successor
- **Link:** [arXiv](#) | [Code](#)

74. **Direct Preference Optimization** | Rafailov et al. | 2023^[14]

- **Training:** Simplified RLHF without separate reward model
- **Citations:** 1,000+ | **Impact:** More efficient alignment training
- **Link:** [arXiv](#) | [Code](#)

75. **Tree of Thoughts: Deliberate Problem Solving** | Yao et al. | 2023^[7]

- **Reasoning:** Structured exploration for complex problems
- **Citations:** 800+ | **Impact:** Enhanced LLM reasoning capabilities
- **Link:** [arXiv](#) | [Code](#)

76. **Retrieval-Augmented Generation for Knowledge-Intensive NLP** | Lewis et al. | 2020^[17]

- **Architecture:** Combining retrieval with generation
- **Citations:** 4,000+ | **Impact:** Foundation for modern RAG systems
- **Link:** [arXiv](#) | [Code](#)

77. **Chain-of-Thought Prompting Elicits Reasoning** | Wei et al. | 2022^[7]

- **Method:** Step-by-step reasoning in language models
- **Citations:** 3,000+ | **Impact:** Improved LLM problem-solving
- **Link:** [arXiv](#)

I ESSENTIAL READING STRATEGY

Priority Levels:

- **MUST-KNOW (Papers 1-15):** Essential for any AI practitioner
- **FOUNDATIONS (Papers 16-25):** Historical context and theory
- **DOMAIN SPECIFIC (Papers 26-72):** Pick your specialization
- **CUTTING EDGE (Papers 73-77):** Stay current with latest developments

Reading Roadmap:

1. **Week 1-2:** Read all 15 MUST-KNOW papers
2. **Week 3-4:** Foundations (16-25) for historical context
3. **Month 2:** Deep dive into your domain of interest
4. **Month 3:** Explore other domains for breadth
5. **Ongoing:** Stay updated with recent breakthroughs

Implementation Strategy:

- **For each paper:** Read, understand, implement key concepts
- **Build portfolio:** 3-5 implementations from different domains
- **Join communities:** Follow authors on Twitter, join Discord servers
- **Stay current:** Follow ArXiv, attend conferences (NeurIPS, ICML, ICLR)

Success Metrics:

- **Beginner → Intermediate:** Understand and can explain all 15 must-knows
- **Intermediate → Advanced:** Deep knowledge in 1-2 domains, broad knowledge across all
- **Advanced → Expert:** Contributing to field, citing and building on these foundations

This list represents the **canonical knowledge** every serious AI practitioner should possess. These papers have shaped the field and continue to influence current research and development.

**

1. https://www.reddit.com/r/MachineLearning/comments/zetvmd/d_if_you_had_to_pick_1020_significant_papers_that/
2. <https://www.nature.com/articles/nature14236>
3. <https://arxiv.org/abs/1412.6980>
4. https://github.com/RPC2/DRL_paper_summary
5. <https://viso.ai/deep-learning/computer-vision-papers/>
6. <https://blog.researchpal.co/academic-research/the-10-most-influential-ai-research-papers-of-all-time/>

7. <https://www.zeta-alpha.com/post/analyzing-the-homerun-year-for-l1ms-the-top-100-most-cited-ai-papers-in-2023-with-all-medals-for-o>
8. <https://papers.baulab.info>
9. <https://www.doradolist.com/papers/21-most-cited-machine-learning-papers>
10. https://www.reddit.com/r/MachineLearning/comments/16ij18f/d_the_ml_papers_that_rocked_our_world_20202023/
11. https://www.reddit.com/r/learnmachinelearning/comments/mm8a4j/mustread_papers_on_generative_a_dversarial/
12. <https://ceur-ws.org/Vol-3742/paper17.pdf>
13. <https://arxiv.org/abs/2312.00752>
14. <https://github.com/SarahRastegar/Best-Papers-Top-Venues>
15. <https://www.linkedin.com/pulse/top-10-breakthrough-computer-vision-technologies-shaping-2024-v1igf>
16. https://en.wikipedia.org/wiki/Deep_learning
17. <https://machinelearningmastery.com/5-breakthrough-machine-learning-research-papers-already-in-2025/>
18. <https://github.com/floodsung/Deep-Learning-Papers-Reading-Roadmap>
19. <https://machinelearningmastery.com/5-most-influential-machine-learning-papers-2024/>
20. <https://www.dataversity.net/brief-history-deep-learning/>
21. https://cvpr.thecvf.com/Conferences/2025/News/Awards_Press
22. <https://github.com/terryum/awesome-deep-learning-papers>
23. <https://www.sciencedirect.com/science/article/pii/S1568494625006891>
24. <https://github.com/youssefHosni/Weekly-Top-Computer-Vision-Papers>
25. <https://www.techrxiv.org/users/778163/articles/912792-deep-learning-historical-overview-from-inception-to-actualization-models-applications-and-future-trends>
26. <https://huggingface.co/papers/trending>
27. https://www.reddit.com/r/MachineLearning/comments/1i39iuh/d_recommendations_of_noteworthy_ai_papers_for/
28. <https://digitalcommons.wku.edu/cgi/viewcontent.cgi?article=4558&context=theses>
29. [http://dtransposed.github.io/blog/2018/12/31/Best-of-GANs-2018-\(Part-1-out-of-2\)/](http://dtransposed.github.io/blog/2018/12/31/Best-of-GANs-2018-(Part-1-out-of-2)/)
30. https://www.mit.edu/~jgabbard/assets/18085_Project_final.pdf
31. <https://neptune.ai/blog/6-gan-architectures>
32. https://docs.pytorch.org/tutorials/intermediate/reinforcement_q_learning.html
33. <https://arxiv.org/pdf/1412.6980.pdf>
34. <https://pmc.ncbi.nlm.nih.gov/articles/PMC8921387/>
35. <https://arxiv.org/abs/1901.00137>
36. <https://proceedings.neurips.cc/paper/2020/file/f3f27a324736617f20abbf2ffd806f6d-Paper.pdf>
37. <https://www.sciencedirect.com/science/article/pii/S2667096820300045>
38. <https://github.com/nightrome/really-awesome-gan>

39. https://www.reddit.com/r/reinforcementlearning/comments/1is773d/must_read_papers_for_reinforcement_learning/
40. <https://www.ewadirect.com/proceedings/ace/article/view/2366>
41. <https://arxiv.org/vc/arxiv/papers/2005/2005.00065v1.pdf>
42. <https://spinningup.openai.com/en/latest/spinningup/keypapers.html>